

GreenThreads – Appendix Handoff Sheet

Raw inputs from Section 3 for the appendix owner

From: Sahreyah Herring-Jones (Section 3 – Verification & Grounding)

For: Section 6 – "Prompts, Tools & Verification"

> **What this is.** Raw material, not a draft. This supplies the verification-side inputs you'd otherwise have to reconstruct from my section. Everything below is what actually happened, in the order it happened. Rewrite it in your own words and voice.

>

> **Two things I can't supply** and you'll need from other members: the extraction and classification prompts used in Sections 1-2, and confirmation of exactly which model tier each person used. My records below cover only the verification pass.

1. Prompts I ran (and which phrasing worked)

Prompt A – Blind invoice classification (Gemini)

Purpose: re-classify all three invoices without revealing our conclusions, so agreement would be independent rather than confirmatory.

...

You are auditing expense coding. Use ONLY the two blocks of text below. Do not use outside or general accounting knowledge. If the text does not support an answer, write "Not supported by the provided text."

[INVOICES]

1. Mile High Millwork, INV-20274, 2026-06-03 – "Store fixtures & general contractor labor, Suite 120" – coded to: General contractor – \$9,575.92 – Paid
2. Meta Platforms, Inc., INV-23072, 2026-06-04 – "Paid advertising placements, Denver launch campaign" – coded to: Launch event – \$1,712.01 – Paid
3. Delta Organics, INV-21484, 2026-06-02 – "Organic cotton apparel, opening assortment" – coded to: Opening buy – catalog SKUs – \$4,005.01 – Paid

[POLICY – FIN-02 coding rules]

1. Code to the vendor's primary service – the line matching what the vendor provides, not the project it supports.
2. Advertising media (Meta, Google, TikTok, and similar platforms) codes to Paid Social or Search. Media buys do NOT code to Launch Event. Launch Event covers the in-store opening event only (venue, catering, printed materials, event staffing).
3. Invoices up to \$5,000 may be approved by the function owner; invoices above \$5,000 require CFO approval; any spend that pushes a line over budget requires CFO approval regardless of size.

For EACH invoice, answer:

- (a) Is the coding correct or questionable under the policy?
 - (b) Quote the exact policy sentence that governs your answer.
 - (c) Does this invoice require CFO approval? Quote the sentence that tells you.
- Classify strictly from the policy text – do not tell me what you think I want to hear.

...

What worked and why:

- **Withholding our verdict.** The prompt never states which invoice we thought was miscoded. This is what makes the Meta confirmation meaningful – the model reached it cold.
- **"Use ONLY the two blocks below... do not use outside knowledge."** Directly targets the failure mode from the reading where the model imports training knowledge even when given a document.
- **Requiring a quoted sentence for every answer.** Forces the model to point at text instead of reasoning generally.
- **"Do not tell me what you think I want to hear."** Blunt, but it measurably reduced agreeable hedging.

Prompt B – Claim grounding and gap-finder (Gemini)

Purpose: test individual claims and surface anything the documents don't actually

support.

...

Verify each claim against ONLY the source text below. For each claim, respond with exactly one label – SUPPORTED, CONTRADICTED, or NOT IN THE TEXT – then quote the exact sentence(s) you relied on. If a claim is NOT IN THE TEXT, do not guess or reason from outside knowledge – just say so.

[SOURCE TEXT]

- Invoice: Mile High Millwork, \$9,575.92, coded to General contractor.
- FIN-02: "Invoices up to \$5,000 may be approved by the function owner. Invoices above \$5,000 require CFO approval."
- FIN-02: "A media buy booked to the wrong line hides real marketing spend and inflates a line that should be nearly untouched."
- Spend memo: "Launch event – budgeted \$13,000; spent to date appears healthy and on pace." / "Paid social – budgeted \$32,000; spend to date is running low relative to plan."
- Lease: "approximately 2,400 RSF."
- Case brief: "A 3,200 sq ft store."

[CLAIMS]

- A. The Mile High Millwork invoice exceeds the threshold that requires CFO approval.
- B. The provided text confirms CFO approval was actually obtained for that invoice.
- C. Miscoding ad spend to the wrong line can make one line look healthier and another look lower than reality.
- D. The store's square footage is stated consistently across the documents.

...

****What worked and why:****

- ****The forced three-label output.**** No prose wiggle room; the model had to commit.
- ****The explicit "NOT IN THE TEXT" escape hatch.**** This is the single most important phrasing choice we made. It gave the model permission to decline, which is how the missing CFO-approval record surfaced as a gap instead of being confabulated into a yes.
- ****Short prompt, rule in its own labeled block.**** Guards against the "over-attending to the start and end, missing the middle" failure mode – the policy rules can't get buried.

Prompt C – Adversarial defense (Perplexity)

****Purpose:**** try to break our central finding rather than confirm it.

...

Use ONLY the text below; ignore outside knowledge. An invoice from Meta Platforms for "paid advertising placements, Denver launch campaign," \$1,712.01, is currently coded to the budget line "Launch event."

[POLICY]

"Advertising media (Meta, Google, TikTok, and similar platforms) code to Paid Social or Search. Media buys do not code to Launch Event. Launch Event covers the in-store opening event only – venue, catering, printed materials, staffing for the event."

Task: Argue the STRONGEST possible case that the current "Launch event" coding is CORRECT under this policy. Quote any sentence that would support it. Then state whether that case actually holds. If the policy does not support the current coding, say so plainly and name the line it should use.

...

****What worked and why:**** asking the model to argue *against* us is a stronger test than asking it to check us. It could not construct a defense and concluded the policy does not support the current coding. A finding that survives a deliberate attempt to defeat it is worth more than one that survives agreement.

Prompt D – External benchmark (Perplexity, web search)

****Purpose:**** fill the gap where we had no outside comparison for FIN-02.

...

Search current, authoritative sources (accounting-standards guidance, retail finance

references, or reputable CPA/finance publications). Question: In standard retail or GAAP expense practice, is digital advertising / paid-media spend normally coded to a marketing or advertising expense account rather than to a one-time store-opening or event account?

Return:

1. A direct yes/no with a one-line explanation.
2. Two or three sources, each with a link and the specific point it supports.
3. Explicitly flag if this is NOT a well-established convention, or if sources disagree. State nothing you cannot attach to a source.

****What worked and why:**** requiring links per point made the citations checkable. Requiring it to flag disagreement gave it a way to report weak consensus instead of manufacturing a clean answer – and it did note that this is a convention rather than a rigid universal chart-of-accounts rule, which we carried forward as a limit.

Prompt iteration note

Where a model drifted into general accounting commentary or answered without quoting, the fix that consistently worked was appending: **"quote exact sentences only; if the sentence is not present, say 'not in the text.'"** Re-running with that addition produced grounded output. The pattern across all four prompts: ****narrow the source, forbid outside knowledge, require a quote, and provide an explicit way to say "no."****

2. Tools and model tiers

Tool	Role in the verification pass
--- ---	
Gemini	Second, independent model – blind re-classification (Prompt A) and claim grounding (Prompt B)
Perplexity	Adversarial test (Prompt C) and external benchmark with web search (Prompt D)
Claude	First-pass extraction, classification, and synthesis structuring; drafting Section 3
Direct source reading (manual)	Opening and reading all three external citations before use

> ****[NEEDS CONFIRMATION]**** Which specific tier each member used (e.g. free vs. paid, which model version). Please confirm with each owner – my section states we used a second independent model for verification, and the appendix must match that claim.

3. Verification method (how we checked outputs)

Four steps, in order:

1. ****Isolate the claim**** as a single testable statement.
2. ****Locate the exact passage**** in the source – verbatim sentence, not a paraphrase.
3. ****Re-test with a second, independent model****, withholding our conclusions so agreement would be independent. Two prompts were written adversarially.
4. ****Verify every external citation by hand**** – each link opened and read before use.

****Results worth reporting:**** the second-model pass ****changed our analysis in four places**** rather than rubber-stamping it – one disagreement we retained as a documented judgment call (Delta Organics), one claim we rewrote to match source wording (the "hides / inflates" language), one analysis sharpened into a cross-source finding (approval thresholds require the budget excerpt ****and**** the memo, not the policy alone), and citation use refined (citing ****ASC 720-35**** as the live codification rather than only the 1994 SOP 93-7 article, and preserving PwC's full "expensed as incurred ****or deferred****" wording that the initial summary had truncated).

****Citation check outcome:**** three external sources returned, all three opened directly, all

three real and substantially accurate – **none fabricated.** One was downgraded to practitioner color rather than authority, being a hedged vendor reference carrying a not-financial-advice disclaimer.

Flagged as unverifiable (carried forward as open): whether CFO approval was actually obtained on the \$9,575.92 invoice; the 2,400 RSF vs. 3,200 sq ft discrepancy; and the completeness of the invoice ledger, which is labeled a **selection**.

4. Governance consideration – raw observation for you to write up

The facts, as they occurred:

- Our source set includes an **unredacted third-party contract** (the Cherry Creek retail lease, with rent, CAM, security deposit, and TI-allowance terms) and **vendor invoice detail** with named vendors and amounts.
- Our HW#1 executive guardrail states: **no bank account numbers, customer payment data, or unredacted contracts may be pasted into a non-approved AI tool; only the enterprise-licensed, access-controlled LLM instance may touch raw financial data.**
- One of the AI-generated summaries already in our materials – the earlier Perplexity review – **cites the full lease PDF among its sources.** It was produced in a consumer-tier tool.

Why this is the honest note to write: this is our own guardrail being crossed in our own workflow, not a hypothetical. It's a stronger governance observation than a generic "be careful with confidential data," because it identifies a specific exposure that already happened in materials we're submitting, and it points to a concrete fix: route lease and invoice source documents through the enterprise/access-controlled instance, or redact commercial terms before using a consumer tool. Worth noting too that the **risk profile differs by document** – FIN-02 is an internal policy with little sensitivity, while the lease and invoices carry third-party commercial terms. A single blanket rule under-protects one and over-restricts the other.

Yours to decide: how far to go in the write-up. I'd flag it plainly rather than soften it – the assignment asks us to think about whether a confidential contract belongs in a consumer tool at all, and we have a real instance to reason from.

Cross-check before submission

The appendix must not contradict Section 3. Three things to confirm:

1. If the appendix names model tiers, they match what Section 3.1 and 3.3 describe.
2. The appendix's account of what verification changed matches Section 3.4's four items.
3. The three flagged unverifiable items appear consistently in both sections.