

GreenThreads Section 3

Verification & Grounding · AI.205 HW#2

Section owner: Sahreyah Herring-Jones · **Feeds:** Section 4 (Synthesis), Section 6 (Appendix)

3.1 Method

Every finding in this brief was treated as unproven until a specific sentence in a source document could be produced for it. Our procedure had four steps:

1. **Isolate the claim.** Each finding was written as a single testable statement, so it could be confirmed or denied rather than partly agreed with.
2. **Locate the passage.** We required the exact sentence from the source – not a paraphrase, not a summary – and recorded it verbatim.
3. **Re-test with a second, independent model.** Findings were re-run through Gemini using prompts that withheld our conclusions, so the model reached its own verdict rather than confirming ours. Two prompts were written adversarially, asking the model to argue *against* our position.
4. **Verify all external citations by hand.** Every source returned by Perplexity was opened directly and read before use.

Two design choices are worth stating, because they follow the course readings on how these models behave. First, each verification prompt included an explicit escape hatch – *"if the text does not support an answer, say 'not supported'"* – which gave the model permission to decline instead of inventing support. This is what surfaced our approval-trail gap rather than papering over it. Second, we kept prompts short and placed the governing policy in its own labeled block, to reduce the risk of the model over-attending to the beginning and end of a long prompt and missing the rule in the middle.

We also treated the two AI-generated summaries already in our materials (a prior Perplexity review and a classification screenshot) as *drafts requiring audit, not as findings.* Their conclusions were re-derived from source text before any were carried forward.

3.2 Claim-to-Passage Grounding Table

#	Claim	Exact supporting passage	Source	Status
1	The Meta invoice is coded to <i>*Launch event*</i>	"Coded to budget line: Launch event . Amount: \$1,712.01 · Status: Paid"	Budget & Invoice Excerpt	Verified
2	Policy requires that spend to code to <i>*Paid social*</i>	"Purchases of advertising media (Meta, Google, TikTok, and similar platforms) code to Paid Social or Search... Media buys do not code to Launch Event."	Policy FIN-02, Rule 2	Verified
3	<i>*Launch event*</i> covers the in-store event only	"Launch Event covers the in-store opening event only – venue, catering, printed materials, staffing for the event."	Policy FIN-02, Rule 2	Verified
4	Miscoding inflates the wrong line and hides real spend from the correct one	"A media buy booked to the wrong line hides real marketing spend and inflates a line that should be nearly untouched."	Policy FIN-02	Verified (wording tightened – see 3.4)
5	The Millwork invoice exceeds the CFO approval threshold	"\$9,575.92" vs. "Invoices above \$5,000 require CFO approval."	Invoice + FIN-02, Rule 3	Verified
6	CFO approval was actually obtained for that invoice	– no passage exists –	–	NOT VERIFIABLE – gap
7	The Millwork coding is correct	"Store fixtures & general contractor labor" coded to <i>*General contractor*</i> ; "Code to the vendor's primary service"	Invoice + FIN-02, Rule 1	Verified
8	The Delta coding is correct	"Organic cotton apparel, opening assortment" coded to <i>*Opening buy – catalog SKUs*</i> (category: Inventory)	Invoice + Budget Excerpt	Verified with stated assumption – see 3.4
9	No budget line is over budget as currently coded	\$9,575.92 of \$78,000; \$1,712.01 of \$13,000; \$4,005.01 of \$29,300	Budget & Invoice Excerpt	Verified by arithmetic

| 10 | The memo reports *Launch event* healthy and *Paid social* low | "Launch event... spent to date appears healthy and on pace." / "Paid social... spend to date is running low relative to plan." | Spend Memo | ****Verified**** (passages); the causal link to #4 is our inference, labeled as such |

| 11 | Contingency reserve is untouched | "the \$20,000 contingency reserve untouched" | Spend Memo | ****Verified**** (self-reported by Finance) |

| 12 | Store square footage is stated consistently | Lease: "approximately 2,400 RSF" / Case Brief: "A 3,200 sq ft store" | Lease + Case Brief | ****CONTRADICTED – discrepancy**** |

| 13 | Paid social has received \$0 of correctly-coded spend | Packet is labeled "Invoices received (****selection****)" | Budget & Invoice Excerpt | ****Partially verified**** – no correctly-coded ad spend is *visible*; the full ledger was not provided |

****Fact vs. judgment.**** Claims 1, 2, 3, 5, 9, 10 and 12 are matters of record – the document says it or it doesn't. Claim 4's *mechanism* is quoted policy, but the conclusion that the Meta invoice ****should be recoded**** is an interpretation: we apply FIN-02 as the controlling standard. We state this plainly rather than presenting a judgment as a finding.

3.3 Second-Model Verification Results (Gemini)

Gemini was given the invoice and policy text with no indication of our conclusions, and asked to classify each invoice and quote the governing rule.

Item	Our finding	Gemini (independent)	Outcome
Meta coding	Questionable	**Questionable** – cited Rule 2 unprompted	Confirmed
Millwork coding	Correct	**Correct** – cited Rule 1	Confirmed
Millwork > \$5,000 → CFO	Yes	**Yes**	Confirmed
CFO approval obtained?	Not in documents	**NOT IN THE TEXT**	Confirmed (gap)
Sq-ft consistency	Contradiction	**CONTRADICTED**	Confirmed
Delta coding	Correct	**Questionable** **Divergence** → 3.4	
Claim 4 as first drafted	Supported	**NOT IN THE TEXT** **Divergence** → 3.4	
Approval sufficiency, small invoices	Function owner may approve	**Cannot determine from policy alone** **Refinement** → 3.4	

****Adversarial test.**** Perplexity was separately asked to build the strongest possible case that the *Launch event* coding is correct under FIN-02. It could not construct one, and concluded the policy does not support the current coding. A finding that survives a deliberate attempt to defeat it is materially stronger than one that survives only agreement.

3.4 Where Verification Changed Our Analysis

Verification is only meaningful if it is capable of altering the result. It altered ours in four places.

****1. Delta Organics – disagreement retained as a documented judgment call.****

Gemini classified this invoice as questionable, reasoning that "Opening buy" names a *purpose* while the vendor's service is *apparel*, and that FIN-02 Rule 1 requires coding to the service, not the project. We considered this and retained our original verdict, for a reason grounded in the budget document: the line is "Opening buy – ****catalog SKUs****" and its stated category is ****Inventory****. Delta supplies apparel; the apparel *is* the catalog SKU inventory. Coding therefore matches what the vendor provides and satisfies Rule 1. The budget excerpt also contains no alternative inventory line to which the invoice could be moved – unlike the Meta invoice, where a correct destination (*Paid social*) exists and is named in the policy.

****Stated assumption:**** we read "Opening buy – catalog SKUs" as the inventory line rather than an opening-event line. If GreenThreads intends that line as event-related, the classification would change.

****2. Claim 4 – wording tightened to match the source.****

As first drafted, the claim said miscoding makes one line "look healthier and another look

lower than reality." Gemini returned NOT IN THE TEXT, because the policy states that miscoding *hides* real marketing spend – it does not say a line "looks lower." The objection was fair, and rather than argue the point we rewrote the claim to track the source language: *inflates the wrong line and hides real spend from the correct one.* The claim is now fully supported by quoted text.

****3. Approval analysis – became a cross-source finding.****

Gemini declined to confirm that function-owner approval was sufficient for the two invoices under \$5,000, noting that Rule 3 has a ****second trigger****: any spend pushing a line over budget requires CFO approval regardless of size, and the policy text alone does not reveal line balances. This is correct. Confirming sufficiency requires the ***budget excerpt*** for the caps and the ***spend memo*** for status: Launch event stands at roughly \$1,712 of \$13,000 and Opening buy at \$4,005 of \$29,300, so neither breaches its cap. ****No single document answers the approval question**** – a point Section 4 should carry, as it demonstrates genuine cross-source dependency rather than parallel reading.

****4. External citations – verified and refined.****

Perplexity returned three sources supporting the view that paid digital media is treated as advertising expense rather than as a store-opening or event cost. All three were opened and read directly. All three were real and substantially accurate; none was fabricated. We refined their use in three ways: we cite ****ASC 720-35**** as the live codification rather than relying only on the 1994 SOP 93-7 article that preceded it; we preserved PwC's full "expensed as incurred ****or deferred****" language, which the initial summary had truncated; and we labeled the vendor-blog source as practitioner color rather than authority, since it is hedged and carries a not-financial-advice disclaimer.

External source	Authority	Use
PwC Viewpoint §6.10, *Advertising costs*	High – curated GAAP guidance	Primary external anchor
CPA Journal, *Reporting on Advertising Costs* (1994, SOP 93-7)	Medium – historical origin of the rule	Shows the convention is long-established
Ramp, *What expense category is social media management?*	Low – vendor reference, hedged	Practitioner color only

****A distinction we are careful not to overstate.**** These sources establish the ****category**** – that paid media is advertising expense, not a capitalizable pre-opening or event cost. They do ****not**** establish the ****line**** – that this invoice belongs on ***Paid social*** specifically. That second step rests on FIN-02, which is GreenThreads' controlling internal standard. Stated in two layers, the conclusion is stronger, not weaker: the Meta coding is inconsistent with general accounting convention ***and*** with the company's own written policy.

3.5 Items We Could Not Verify (Flagged)

Per the assignment's instruction to flag anything unverifiable, three items are carried forward as open rather than resolved:

- **CFO approval on the Millwork invoice (\$9,575.92).**** The invoice exceeds the \$5,000 threshold and therefore required CFO approval. No document in our set records whether that approval was obtained. This is a ****controls gap****, not a finding of wrongdoing – the evidence is simply absent.
- **Store square footage.**** The lease states approximately 2,400 RSF; the case brief states 3,200 sq ft. We cannot resolve this from the documents. Any per-square-foot cost, CAM, or buildout assumption should be treated as unreliable until reconciled.
- **Completeness of the invoice ledger.**** The packet is explicitly a ***selection***. We therefore cannot assert that ***Paid social*** has received no spend – only that none is visible in what was provided.

Handoff Note – Section 4 (Synthesis)

****What is certified and may be asserted plainly:****

- The Meta invoice is coded to *Launch event*; FIN-02 requires *Paid social*; the policy explicitly forbids media buys on *Launch Event*. Confirmed by a second model reaching this independently, and supported externally by ASC 720-35 treatment of advertising costs.
- Millwork and Delta codings are correct (Delta with the stated assumption in 3.4).
- No line is over budget as coded; contingency is untouched.
- The memo's "healthy" and "running low" readings are downstream of the miscoding.

****What must be labeled as interpretation, not fact:****

- That the invoice *should be recoded* – this applies FIN-02 as the controlling standard; the invoice itself says only where it was booked.
- The causal link between the miscoding and the memo's conclusions.

****What must be presented as a gap, not a conclusion:****

- The missing CFO approval record. Say "not evidenced," never "not approved."
- The square-footage discrepancy.
- The partial ledger, which limits any claim about *Paid social* at \$0.

****Wording to use verbatim, because it is what the source supports:**** miscoding *"inflates"* the wrong line and *"hides"* real spend from the correct one.

****The strongest available synthesis point:**** FIN-02's own rationale predicted this exact failure before it occurred – the policy warned that a media buy on the wrong line hides real marketing spend and inflates a line that should be nearly untouched, and that is precisely what the memo now reports. The control was designed correctly and not executed. That is a *controls-execution* gap, not a controls-design gap, and it is the sharpest thing our documents support.